

Term-Indexing for the Beagle Theorem Prover

Timothy Clarence Richard Cosgrove

A subthesis submitted in partial fulfillment of the degree of
Bachelor of Science (Honours) at
The Department of Computer Science
Australian National University

September 2013

© Timothy Clarence Richard Cosgrove

Typeset in Palatino by \TeX and $\text{\LaTeX} 2_{\epsilon}$.

Except where otherwise indicated, this thesis is my own original work.

Timothy Clarence Richard Cosgrove
17 September 2013

For Dana.

Acknowledgements

Thank you to my Supervisor and all...

Abstract

This should be the abstract to your thesis. . .

x

Contents

Acknowledgements	vii
Abstract	ix
1 Introduction	1
1.1 Motivation	1
1.1.1 A Theoretical Framework	1
2 Background	3
2.1 First-Order Logic Terms and Notation	3
2.1.1 FOL basics	3
2.1.2 Calculi and FOL problems	3
2.1.3 The Superposition Calculus	3
2.1.4 Specialised Syntax in this Paper	3
2.2 Automated Reasoning and Theorem Proving	4
2.2.1 SPASS	4
2.2.2 Vampire	4
2.2.3 E	4
2.3 Term Indexing	4
Top Symbol Hashing	4
Discriminant Trees	4
2.4 Fingerprint Indexing	4
2.5 The <i>Beagle</i> Theorem Prover	5
2.5.1 The Heirachic Superposition with Weak Abstraction Calculus . .	5
2.5.2 <i>Beagle</i> 's Shortcomings	5
2.6 Tools Used	5
2.6.1 Scala	5
2.6.2 VisualVM	5
2.6.3 Eclipse	5
3 Implementing Fingerprint Indexing	7
3.1 Structure of <i>Beagle</i>	7
3.1.1 Syntax and Data Structures	7
3.1.2 Main Program Loop	7
3.2 Building the Fingerprint Indexer	8
3.2.1 Objects and Datatypes	9
Fingerprint Features	9

Positions	9
Term Fingerprints	9
Fingerprint Index	9
3.2.2 Building Term Fingerprints	9
3.2.3 Adding Terms to the Index	10
3.2.4 Retrieving Compatible terms	10
3.2.5 Matching with Subterms	12
Data loss	12
Non-unique Representation	12
3.2.6 Term Tracing	12
3.2.7 Unit Testing with ScalaTest	13
3.3 Adding Indexing to <i>Beagle</i>	13
3.3.1 Initial Problems	13
3.3.2 Refactoring Current Implementation	13
3.4 Extending the Indexer	13
3.4.1 Matching and Simplification in <i>Beagle</i>	13
Negative Unit Simplification	14
Demodulation	14
3.4.2 First Attempt at Indexing Simplification	14
3.4.3 Generalising our FingerprintIndex	14
3.4.4 Applying new Indices to Simplification	14
3.5 Tailoring to <i>Beagle's</i> Heirachic Superposition with Weak Abstraction Calculus	14
3.5.1 Extending the Unification Table with Term Layers	14
3.5.2 Extended Matching Table	16
4 Results	17
4.1 Beagle Before Indexing	17
4.1.1 Points of Improvement	17
4.2 Indexing Metrics	17
4.2.1 Problem Selection	17
4.2.2 Speed	17
4.2.3 False Positives	17
4.3 Indexing Subsumption	17
4.3.1 Metric Results	17
4.3.2 Comparison	18
4.4 Indexing Simplification and Matching	18
4.4.1 Further Intstrumentation	18
4.4.2 <i>Beagle</i> with Simplification Improvements	18
4.5 Tailored Improvements	18
4.5.1 Layer Checking	18
4.5.2 Metric Results	18

5 Conclusion	19
5.1 Why this is a Very Clever Thesis	19
5.2 Future Improvements	19
5.2.1 Extended Data Structures	19
5.2.2 More Fingerprint Indices	19
5.2.3 Extensions to Fingerprint Indexing	19
Symbol count	19
Dynamic Fingerprinting	19
5.2.4 Combining Indexing Techniques	19
5.3 Final Thoughts	19
A Result Tables for TPTP Selection	21
B Full TPTP run?	23
Bibliography	25

Introduction

1.1 Motivation

- Describe beagle
- Advantages of beagle
- drawbacks
- some instrumentation

1.1.1 A Theoretical Framework

Background

2.1 First-Order Logic Terms and Notation

This thesis focuses around the extension of *beagle*, a *first-order logic* (FOL) theorem prover. In order to understand beagle's purpose and functions a basic understanding of the FOL logical system is required. This section provides a rudimentary overview of FOL syntax and uses; but also includes an explanation of any specialised terms and notation used throughout the paper.

2.1.1 FOL basics

2.1.2 Calculi and FOL problems

2.1.3 The Superposition Calculus

Should contain

- Variables
- Symbols
- Predicates
- Quantifiers
- Notion of soundness and completeness
- Description of a 'calculus'

2.1.4 Specialised Syntax in this Paper

- Positions

2.2 Automated Reasoning and Theorem Proving

Automated Reasoning is a rapidly growing field of research where computer programs are used to solve problems stated in first order logic statments or other formal logics.

Some existing theorem provers include:

2.2.1 SPASS

[Weidenbach et al. 1999]

2.2.2 Vampire

[Riazanov and Voronkov 1999]

2.2.3 E

[Schulz 2002]

Should contain

- Theorem prover examples

2.3 Term Indexing

Term indexing is a technique used to better locate logical terms which match rules in a prover's calculus.

Top Symbol Hashing

Discriminant Trees

2.4 Fingerprint Indexing

Fingerprint Indexing is a recent technique developed by Schulz [2012], the creator of the E prover.

Table 2.1: Fingerprint matches for Unification [Schulz 2012, p6]

	f_1	f_2	A	B	N
f_1	Y	N	Y	Y	N
f_2	N	Y	Y	Y	N
A	Y	Y	Y	Y	N
B	Y	Y	Y	Y	Y
N	N	N	N	Y	Y

Table 2.2: Fingerprint matches for Matching [Schulz 2012, p6]

	f_1	f_2	A	B	N
f_1	Y	N	N	N	N
f_2	N	Y	N	N	N
A	Y	Y	Y	N	N
B	Y	Y	Y	Y	Y
N	N	N	N	N	Y

2.5 The *Beagle* Theorem Prover

The core implementation of *beagle* was developed by Peter Baumgartner et al. of NICTA. Its purpose was to demonstrate the capabilities of the *Weak Abstraction with Heirachic Superposition Calculus*; which allows the incorporation of prior knowledge via ‘background reasoning’ modules.

2.5.1 The Heirachic Superposition with Weak Abstraction Calculus

2.5.2 *Beagle*’s Shortcomings

2.6 Tools Used

2.6.1 Scala

As mentioned above *beagle* is written in *Scala*, the Scalable Language. Scala is a functional language and may be confusing to those who are not familiar with the functional programming paradigm. This thesis will contain occasional snippets of Scala code; but note that any snippets used will be accompanied by an explanation and in general an understanding of Scala/functional programming is not required.

2.6.2 VisualVM

2.6.3 Eclipse

Implementing Fingerprint Indexing

3.1 Structure of *Beagle*

Making any extension to the *beagle* project (or any sizeable project for that matter) will obviously require a solid understanding of the existing codebase. This section provides an overview of any existing Scala classes and their structure which is relevant to the implementation of the Fingerprint Index.

3.1.1 Syntax and Data Structures

3.1.2 Main Program Loop




Figure 3.1: Class structure for internal representation of logical formulae.

3.2 Building the Fingerprint Indexer

The first step in adding Fingerprint Indexing to *beagle* is creating the indexer itself; a Scala class which will manage the index and provide functions for adding to it and retrieving from it.

3.2.1 Objects and Datatypes

Fingerprint Features

```

1  /** Pseudo enumerated type for fingerprint features */
2  sealed abstract class FPFeature
3  case object FPA extends FPFeature
4  case object FPB extends FPFeature
5  case object FPN extends FPFeature
6  case class FPF(val subterm : FunTerm) extends FPFeature

```

Listing 1: Data type for the 4 Fingerprint Features [Schulz 2012, p5]

Positions

We implement positions in the simple/naïve manner, simply as a list of Integers (where Nil is used to index the top-level term).

Term Fingerprints

Fingerprint Index

The final data structure we require is the

```

1  /** Algebraic Data type for our index. Either we are at a leaf
2   * (set of terms) or must continue traversal via the map. */
3  sealed abstract class Index
4  case class Leaf(set: Set[Term]) extends Index
5  case class Node(map: HashMap[FPFeature, Index]) extends Index

```

Listing 2: Data type for the actual term index. [Schulz 2012, p7]

Note that this Index object does not necessarily take up a significant portion of memory. All Terms are already stored within the ClauseSet object (See Figure 3.1); so the Index itself will generally only add a fairly lightweight structure of pointers.

3.2.2 Building Term Fingerprints

The following block of Scala code extracts a single Fingerprint Feature from a Term at the given position.

```

1  /** Extract the operator at position pos. Note that matching Var
2    * and Funterm is an exhaustive pattern for Term. */
3  def extractFeature(term: Term, pos: Position) : FPFeature = pos match {
4    case p :: ps => term match {
5      case t:FunTerm => try {extractFeature(t.args(p), ps) }
6                        //Non-existent position, return N
7                        catch {case e:IndexOutOfBoundsException => FPN}
8      // Found variable BEFORE end of position, return B
9      case t:Var      => FPB
10   }
11   // Reached end of position, check symbol
12   case Nil          => term match {
13     case t:FunTerm => FPF(t.op) // Found function symbol, return it
14     case t:Var      => FPA      // Found variable, return A
15   }
16 }

```

Listing 3: Scala code to extract fingerprint features for matching.

3.2.3 Adding Terms to the Index

3.2.4 Retrieving Compatible terms

Our Index framework is now capable of creating Fingerprints and storing Terms in our pointer structure. The next task in building our index is allowing retrieval of terms.

To compare two fingerprints with each other we look at them side-by-side and check that each position shows a Y in the Fingerprint unification table.

```

1  /** Check two Fingerprint features for compatibility based
2   * on the unification table (See page 6 of [Shulz 2012]).
3   * This table is reduced to 4 cases:
4   * - True if Features are equal,
5   * - True if at least one Feature is B,
6   * - True if at least one Feature is A; but no Ns,
7   * - False otherwise */
8  def compareFeaturesForUnification
9      (a:FPFeature, b:FPFeature) : Boolean =
10      (a == b) ||
11      (Set(a,b) match {
12          case x if (x contains FPB) => true
13          case x if (x contains FPA) => !(x contains FPN)
14          case _ => false})

```

Listing 4: Scala implementation of the Fingerprint unification table. [Schulz 2012, p6]

To check whether or not two fingerprints match is now a simple matter of iterating through the list and checking that each position is a match according to our unification table check. This side-by-side comparison could potentially be improved by employing some sort of hashing function;

```

1  def retrieveCompatible
2      (t: Term, fp: Fingerprint, index: Index) : TermSet =
3      (fp, index) match {
4          //Collect all compatible (Feature,Index) pairs and continue traversal
5          case (fp::fps, Node(map)) => (for ((k,v) <- map if compare(fp, k))
6              yield retrieveCompatible(t, fps, v))
7          //Collapse all retrieved sets together with the union operator (:::)
8          .foldLeft( Nil:TermSet ) ((a,b) => a ::: b)
9          case (Nil, Leaf(set)) => set
10         case (_, Node(_)) => throw new IllegalArgumentException
11             ("Fingerprint is over but we are not at a leaf")
12         case (_, Leaf(_)) => throw new IllegalArgumentException
13             ("Reached a leaf but Fingerprint is not over")
14     }

```

Listing 5: Scala code to collect compatible terms from the index.

3.2.5 Matching with Subterms

In Shulz’s paper the indexer he describes is only interested in the unification of full, top-level terms with other top-level terms [Schulz 2012]. Recall however the main superposition rules for *beagle*’s resolution calculus [Baumgartner and Waldmann 2013]:

Notice in particular the condition that *s* may match against a *subterm* of *l*. This does not match the usual required definition of unification; where we only match top-level terms.

Thus our Fingerprint Index must be able to collect all possible matches for *subterms*. To do this we will extend *beagle*’s Term object with a function collecting all subterms; along with the position they were extracted from. For variables and constants this is trivial; we just return the symbol and Nil for the subterm position.

```

1  /** Retrieve all subterms along with their position */
2  def subtermsWithPos : List[(Term, List[Int])] =
3    (thisterm, Nil) :: (for
4      ((arg,      argpos)      <- args.zip(args.indices;
5       (subarg, subargpos) <- arg.subtermsWithPos)
6       yield (subarg, argpos::subargpos))

```

Listing 6: Recursively grab all subterms from a complex term.

So our indexer will now be capable of comprehensively indexing all subterms as required. However this ability has been introduced at the cost of a couple of grave errors which, if not detected at this stage, would have destroyed the Fingerprint Index altogether.

Data loss

In listing 2 we presented the actual data structure for storing indexed Terms. It presented the leaves of this index as simply a set of term objects; but this is actually no longer sufficient for our current uses. Now that we are indexing full subterms we must be able to extract the original parent term; which is impossible given only the low-level subterm representation.

Non-unique Representation

It is possible to fix our data loss issue by giving each Expression a pointer to its parent Expression (see Figure 3.1). However, as our index currently only keeps a Set of Terms, If multiple terms possessing the same subterm are ever indexed we will only keep one of them; losing any representation of the other!

3.2.6 Term Tracing

All the above issues can be solved by adding *Term Traces*.

Term tracing also solves a horrendous error introduced by matching subterms.

By indexing subterms we lose information. Must construct a trace in order to restore all required data

Note that adding subterms and their traces has resulted in a significant increase in the size of our Index object. This is negligible however since memory is cheap; and we are generally only concerned with speed.

3.2.7 Unit Testing with ScalaTest

As with any component of a large software project; it is vital to ensure that the Fingerprint Index functions on its own. Otherwise if there are issues when adding the indexer to functional use there will be no way of knowing what component is causing the problem.

3.3 Adding Indexing to *Beagle*

We now have a class fully capable

$$\frac{l \approx r \vee C \quad s[u] \approx t \vee D}{\text{abstr}((s[r] \approx t \vee C \vee D)\sigma)}$$

3.3.1 Initial Problems

3.3.2 Refactoring Current Implementation

Actually making use of our indexer class will require significant modification to *beagle*'s structure and proving sequence.

Refer to class and main loop flow diagrams from 3.2

3.4 Extending the Indexer

Beagle is now fully capable of indexing its terms for superposition. There are still many ways in which this form of indexing may be improved; but it is likely to be far more effective to re-examine where *beagle* now spends most of its runtime and look into other areas which may be improved. By instrumenting the (now indexed) *beagle* in VisualVM (see Section 4.1.1 and 4.4.1) we observe that the most significant runtime cost is now most often in simplification.

3.4.1 Matching and Simplification in *Beagle*

Beagle has several simplification rules to aid the logical inference process. Generally these rules are not part of the actual rule based calculus but rather just implement some external insights which allow some special cases to be pre-processed

Negative Unit Simplification

$$\frac{a \not\approx b \quad a \approx b \vee C}{C}$$

Demodulation

3.4.2 First Attempt at Indexing Simplification

Re-using current index produced little improvement. Cost of indexing subterms, matching against equations with *\$equal*

3.4.3 Generalising our FingerprintIndex

explain available options, their reasoning and their implementations.

3.4.4 Applying new Indices to Simplification

3.5 Tailoring to *Beagle's* Heirachic Superposition with Weak Abstraction Calculus

In this section we discuss the thought process in developing and implementing extensions to Fingerprint Indexing in order to better tailor it to *beagle's* rather unique logical calculus.

3.5.1 Extending the Unification Table with Term Layers

In the Heirachic Superposition with Weak Abstraction Calculus all terms have a concept of being 'Foreground' or 'Background'. In Section 2.5 we discussed this concept; referring to it as the *layer* of a term. It is worht noting at this stage that computing the layer of a term is cheap (or rather, zero, as it is computed on the fly during term generation and stored for later use).

Recall the four original fingerprint feature symbols from 2.3:

- *f*: arbitrary constant function symbols.
- **A**: Variable at the exact position.
- **B**: A variable could be expanded to meet the position.
- **N**: Position can never exist regardless of variable assignment.

We introduce two new fingerprint features: **A+** and **B+**. These symbols will be used for the same purpose as the original **A** and **B**, but only for '*background*' or '*abstraction*' variables. These variables can only be used for pure background terms; a fact we may use to restrict the possible matches for unification.

The layeredness of function symbols is also relevant to our comparison. *f+* in the following table signifies a position where the entire subterm from this position

downwards is 'pure background'. Keep in mind that this definition is slightly different to the definition for **A**⁺ and **B**⁺; as we must consider all function symbols below f itself.

At this point it is important to note that these added fingerprint features slightly modify the original **A**, **B** and f features. These features will now only represent the foreground layered positions.

Table 3.1 displays the unification table with our new background feature symbols. The table has grown to be a considerable size. Refer to the original unification table (Table 2.1) for an in-depth explanation of how this table should be interpreted [Schulz 2012].

Table 3.1: Fingerprint matches for unification; extended by considering term layers.

	f_1	f_2	A	B	N	f_1^+	f_2^+	A ⁺	B ⁺
f_1	Y	N	Y	Y	N	N	N	N	N
f_2	N	Y	Y	Y	N	N	N	N	N
A	Y	Y	Y	Y	N	Y	Y	Y	Y
B	Y	Y	Y	Y	Y	Y	Y	Y	Y
N	N	N	N	Y	Y	N	N	N	Y
f_1^+	N	N	Y	Y	N	Y	N	Y	Y
f_2^+	N	N	Y	Y	N	N	Y	Y	Y
A ⁺	N	N	Y	Y	N	Y	Y	Y	Y
B ⁺	N	N	Y	Y	Y	Y	Y	Y	Y

Note that as this table is for unification it is symmetric along the leading diagonal (as in the original unification table); so we need only discuss the lower triangle of the matrix. Furthermore, notice that the bottom right segment of the table is actually identical to the original unification table. This is expected as when we compare two pure background features the comparison behaves normally.

We will justify the new section of the table line by line:

- Background function symbols (f^+): Recall that this feature is only applicable if the entire subterm below f is pure background. Therefore it does not match the foreground version of the same symbol. It does however match both **A** and **B**. This is required since these symbols still match 'impure' background variables; which may be expanded to either foreground or pure background terms.
- Abstraction variables (**A**⁺): Similarly to the pure background function symbol feature, this feature cannot match any terms which sit in the foreground. It can however match both **A** and **B** as they may represent either foreground or background expressions.
- Potential expansion of an abstraction variable (**B**⁺): Same as for **A**⁺ but can also match **N**.

To go with this table we present its corresponding Scala matching code in Listing 7. Unfortunately the steep increase in table size results in the amount of code required exploding. It also becomes impossible to use our earlier trick of Set matching; due to the need for parametrised Fingerprint symbols (i.e. **A**+ and **B**+ represented as **FPA**(true) and **FPB**(true)).

```

1  /** Check two Fingerprint features for compatibility based
2    * on the *extended* unification table (See table in report).*/
3  def compareFeaturesForUnification
4    (a:FPFeature, b:FPFeature) : Boolean =
5    (a,b) match {
6      case (FPF(f1), FPF(f2))    => (f1.op == f2.op) &&
7                                   (if (f1.isFG || f2.isFG)
8                                     (!f1.isPureBG && !f2.isPureBG)
9                                     else true)
10     case (FPF(f), FPB(true)) => f.isPureBG
11     case (FPB(true), FPF(f)) => f.isPureBG
12     case (_, FPB(_))         => true
13     case (FPB(_), _)         => true
14     case (FPF(f), FPA(true)) => f.isPureBG
15     case (FPA(true), FPF(f)) => f.isPureBG
16     case (FPN, FPA(_))       => false
17     case (FPA(_), FPN)       => false
18     case (_, FPA(_))         => true
19     case (FPA(_), _)         => true
20     case (FPN, FPN)          => true
21     case _                   => false
22   }

```

Listing 7: Scala code to extract fingerprint features for extended layer matching.

3.5.2 Extended Matching Table

Results

4.1 Beagle Before Indexing

We have stated several times throughout this report that the key area of improvement for *beagle* is in clause resolution. However we have yet to provide any evidence to that fact. Here we provide some results from initial investigations (before Fingerprint Indexing was implemented) used to identify the key areas of improvement for *beagle*.

4.1.1 Points of Improvement

Refer to results for instrumentation; showing what may be improved with indexing.

4.2 Indexing Metrics

4.2.1 Problem Selection

Not going to run everything against all of TPTP. Will take out a subset of TPTP (25–50 problems) run them against this set. Show the set and justify inclusions/exclusions

4.2.2 Speed

4.2.3 False Positives

Explain this metric and how it impacts performance

4.3 Indexing Subsumption

4.3.1 Metric Results

We observe many, even after a myriad of optimisations. Notice that this is due to the structure of *beagle*; many retrieved terms are cheaply thrown out due to other conditions (such as being a parent term, being ordered, etc.) and do not significantly impact performance.

Refer to email from Stephan. compare false positive results

4.3.2 Comparison

4.4 Indexing Simplification and Matching

4.4.1 Further Instrumentation

4.4.2 *Beagle* with Simplification Improvements

4.5 Tailored Improvements

4.5.1 Layer Checking

4.5.2 Metric Results

Conclusion

5.1 Why this is a Very Clever Thesis

5.2 Future Improvements

Here we list some thoughts on possible improvements to *beagle* and Fingerprint indexing in general; which were either not relevant to the current work or cut due to time constraints.

5.2.1 Extended Data Structures

The focus of this thesis has been with high-level logical improvements. Add some low-level speed at the cost of memory.

Drop the tree structure. Have an entry for each possible fingerprint and add terms to each bin they are compatible with during indexing. Moves runtime from retrieval (called VERY often) to indexing (called barley ever) but creates memory EXPLOSION

5.2.2 More Fingerprint Indices

5.2.3 Extensions to Fingerprint Indexing

Symbol count

Dynamic Fingerprinting

5.2.4 Combining Indexing Techniques

5.3 Final Thoughts

Result Tables for TPTP Selection

Full TPTP run?

Bibliography

- BAUMGARTNER, P. AND WALDMANN, U. 2013. Hierarchic superposition with weak abstraction. In M. BONACINA Ed., *Automated Deduction – CADE-24*, Volume 7898 of *Lecture Notes in Computer Science*, pp. 39–57. Springer Berlin Heidelberg. (p.12)
- RIAZANOV, A. AND VORONKOV, A. 1999. Vampire. In *Automated Deduction – CADE-16*, Volume 1632 of *Lecture Notes in Computer Science*, pp. 292–296. Springer Berlin Heidelberg. (p.4)
- SCHULZ, S. 2002. E – A Brainiac Theorem Prover. *Journal of AI Communications* 15, 2/3, 111–126. (p.4)
- SCHULZ, S. 2012. Fingerprint indexing for paramodulation and rewriting. In B. GRAMLICH, D. MILLER, AND U. SATTLER Eds., *Automated Reasoning*, Volume 7364 of *Lecture Notes in Computer Science*, pp. 477–483. Springer Berlin Heidelberg. (pp.4, 5, 9, 11, 12, 15)
- WEIDENBACH, C., AFSHORDEL, B., BRAHM, U., COHRS, C., ENGEL, T., KEEN, E., THEOBALT, C., AND TOPIĆ, D. 1999. System description: Spass version 1.0.0. In *Automated Deduction – CADE-16*, Volume 1632 of *Lecture Notes in Computer Science*, pp. 378–382. Springer Berlin Heidelberg. (p.4)